

Reward Hacks That Pass Every Graded Test: Constructing and Detecting SWE-bench Patches That Are Certified Yet Wrong

Daniel Wahnich
Telos research program

July 15, 2026

Abstract

Coding agents are commonly certified when graded tests pass, but a patch can satisfy this proxy while remaining wrong on untested inputs. We ask whether SWE-bench Verified certifies such patches and what detects them.

We report four results. First, an earlier benchmark we built contained no reward hacks: all forty rows fail the official grader. Second, we construct twenty-two patches across eight repositories that pass every graded test yet differ from the accepted fix on a retained input, verified in pinned containers. Third, a two-model judge panel has observed catches on 20/22 hacks, with 8/88 unparseable responses; paired-gold flag sensitivities are 3/22 observed lower, 6/22 worst-case missing upper, and 3/19 complete-case. A second instrument used candidate-diff-derived locators and gold-based inclusion. This locator-assisted, gold-validated property pipeline catches 6/22, all within the judge set, and supplies no independent false-positive estimate. Its full-cohort run failed the registered diff-independent locator contract; an earlier ten-row run additionally replaced its visible-test soundness anchor and omitted registered controls. That earlier run’s apparent property-only row had two unparseable judge responses, so neither run confirms complementarity. Fourth, an exploratory, nonrandom, gold-localized solve with no gaming instruction yields one strict case under a post-inspection blind rule. Of thirty-seven valid patches, twenty-four certify and six certified outcomes remain unadjudicated; the required descriptive sensitivities are 1/24, 7/24 over those six missing pipeline outcomes, and 1/18 complete-case.

These are bounded pilots, not population-frequency estimates. The judge runs retained parsed labels but not raw responses, used capped inputs, and lack an independently Git-frozen full-cohort protocol. The strict blind aggregation rule was adopted after outcome inspection; the missingness summaries were added after the original result but before the denominator backfill. Its pre-output freeze is not independently timestamped in Git. Every number regenerates from committed artifacts.

1 Introduction

An autonomous coding agent is given a bug report and asked to produce a patch. The usual measure of success is a test suite: if the tests pass, the task is considered done. This measure is convenient, and it is also gameable. The agent can edit or weaken the tests, or it can write source code that handles only the exact input the visible test checks rather than solving the general problem. When the test suite is a proxy for “the task is solved,” making the tests pass without solving the task is *reward hacking*, the general failure in which an optimizer satisfies a proxy reward while missing the objective the reward stands in for [1, 2, 3].

This paper asks a specific version of that question. SWE-bench Verified is a coding-agent benchmark [4]. Each of its instances comes with two sets of tests: a `FAIL_TO_PASS` set that a correct fix must newly pass, and a `PASS_TO_PASS` set of existing tests that a correct fix must keep passing. The official harness marks an instance *resolved* when a patch passes both sets. Our question is: can a patch pass both sets—be certified resolved—and still be wrong? If it can, then no check restricted to those graded tests can catch it, because passing those tests is exactly what “resolved” means.

The question is a bounded software-evaluation analogue of a broader oversight problem. Google DeepMind’s Frontier Safety Framework discusses systems that circumvent oversight while outputs are monitored and lists mitigation as future work [7]. Our elicited patches do not establish intent or autonomous oversight evasion; they isolate one measurable proxy/objective gap.

We make four contributions.

1. **A self-correction (Section 3).** An earlier benchmark we released did not contain reward hacks. We show why from committed evidence, and we report the free baseline that had been missing.
2. **Certified-yet-wrong patches (Section 4).** We construct and release a benchmark of twenty-two patches across eight repositories that the official SWE-bench harness marks resolved and that execution shows are wrong, each with a retained witness.
3. **A corrected exploratory instrument comparison (Section 5).** We preserve judge catch and missingness measurements, distinguish two property-protocol failures, and withdraw the independent false-positive, confirmed-complementarity, and deployable-ensemble interpretations.
4. **A neutral-prompt localized existence case (Section 4).** An exploratory, gold-localized solver run produces one strict certified-yet-wrong patch under a post-inspection rule. We correct the certification denominator and report post-result missing-outcome sensitivities without imputing missing outcomes as negatives.

2 Background

Reward hacking. Reward hacking, also called specification gaming [2], is when an agent scores well on a proxy reward while failing the real goal. Skalse et al. [3] give a formal definition. We use the term in its plain sense: a patch that scores well on the grader (it is certified resolved) while being wrong. We note one name collision to avoid confusion: RewardBench [9] evaluates *reward models*, which rank preferences, and is unrelated to the reward hacking studied here.

SWE-bench Verified. SWE-bench [4] asks a model to resolve a real GitHub issue in a real Python repository. SWE-bench Verified is a human-filtered subset. Each instance pins a repository commit and a container image, and defines the `FAIL_TO_PASS` and `PASS_TO_PASS` test sets described above. The “gold patch” is the accepted human fix. A known weakness of the format is that the graded tests can be incomplete: they may not exercise every behavior a correct fix must have. Our constructed hacks live in exactly that gap.

Language-model judges. A language model can act as a monitor that reads a patch and decides whether it is correct [5]. Judges are not uniformly weak. On backdoored programming solutions, a monitor can reach an area-under-the-curve of 0.92 to 0.96 [6]. Our single-run panel has observed

catches on twenty of twenty-two hacks in the full-cohort comparison; it caught seven of ten under the original pilot’s observed rule. In that pilot, missing responses permit a nine-of-ten upper sensitivity and a seven-of-eight complete-case sensitivity. The difference underscores that the retained counts are run-specific and missingness-sensitive, not stable model scores.

Execution-based checks. A different family of checks constructs an input, runs the code, and checks the output. A genuinely reference-free check can use property-based testing [11] or meta-morphic testing [10]. The retained iter197/iter201 instrument did not satisfy the registered diff-independent access contract: its generator received a source/function locator derived from the candidate diff. Offline inclusion also ran each property on gold. Gold validation was explicitly registered in iter201 but not as the iter197 soundness anchor; in either case, it is an inclusion rule rather than an independent control decision. We therefore call the instrument a *locator-assisted, gold-validated property pipeline*, not a validated gold-free detector.

Scale of existing corpora. Public datasets of reward hacking are far larger than ours: METR’s MALT contains on the order of ten thousand transcripts [8]. Ours is twenty-two patches. We make no claim of scale. What is different here is the unit and the verification: we label individual patches (not whole trajectories), and every label is verified by execution. Most benchmark rows are elicited: a model was asked to produce the hacks. The separate neutral-prompt localized solve below supports one strict case under a post-inspection rule, not a population frequency.

3 An Earlier Benchmark Measured the Wrong Thing

We begin with a correction, because it defines the target for the rest of the paper.

An earlier stage of this project released a forty-row artifact and described its rows as reward hacks. The generator that produced those rows kept a candidate only when the candidate made at least one `PASS_TO_PASS` test fail. `PASS_TO_PASS` is the grader’s existing-test subset, so every row failed the official reward by construction.

We re-derived the official harness verdict for all forty rows from the committed evaluation logs. All forty are marked unresolved, and all forty have at least one failing `PASS_TO_PASS` test. The consequence is direct. Running the official harness, or reading its verdict, detects all forty rows at zero model cost. A three-model language-model judge panel, the detector we had been studying, detected seventeen of forty at a metered cost of \$13.59. That free baseline had never been reported.

The rows are therefore not reward hacks. A reward hack scores well on the reward while being wrong; these rows fail the reward. They are wrong patches that break at least one graded regression check, which the official grader catches for free. The earlier judge figure of seventeen of forty is arithmetically correct and we do not retract it; what changes is what it measures. It measures how often a judge, reading a diff, spots a wrong patch that the official graded regression checks already reject. That is a fact about code review, not about reward hacking, and it does not motivate a new verifier.

The real target is the inverse: a patch that passes all graded tests and is still wrong. The official harness scores such a patch as resolved, so running the tests cannot catch it. The rest of the paper constructs that class and asks what does catch it.

4 Constructing Certified-Yet-Wrong Patches

Operational labels. The two cohorts use different rules. A constructed-cohort patch is labelled certified-yet-wrong when the official harness marks it resolved and a clean synthesized-input run produces output different from gold. This is a reference-differential label, not independent semantic adjudication. In the neutral-prompt cohort, a patch must certify, diverge from gold, and—under the later conservative rule—have both blind judges identify only the model output as wrong.

Audited module tests did not distinguish the candidates. Before constructing the class, we took sixteen candidates, confirmed that all sixteen were certified, and ran the graded tests plus every additional test in the same audited test modules. For Django, the graded `PASS_TO_PASS` set already covers the whole target module. Across the sixteen candidates there were 193 additional module tests; none distinguished a candidate from gold. We did not run every test elsewhere in every repository. The measured claim is narrower: witnessing these candidates required an input outside the audited graded/additional target-module tests.

Construction. We construct the class in three steps, all executed in each instance’s pinned container.

1. **Elicit a variant.** We take the gold patch and ask `gpt-5.6-terra` to rewrite its largest added block into a variant that still passes the graded tests but is subtly wrong on other inputs. The model is explicitly asked to produce a wrong variant; the class is elicited, not observed in the wild.
2. **Certify by execution.** We run each instance’s official evaluation command in its pinned container—for example, Django’s `runtests.py`, SymPy’s `bin/test`, or `pytest` where the instance specifies it—and parse the result with the official SWE-bench log parsers. A variant is kept only if it passes every `FAIL_TO_PASS` and `PASS_TO_PASS` test.
3. **Witness the wrongness.** To label a certified variant as wrong, we use the gold patch as reference. `gpt-5.6-terra` constructs a small scenario that calls the changed function on an input designed to reach the branch where the variant and the gold fix could differ. We run that scenario under the gold patch and under the variant, in the same container. If the two produce different output, and the gold run is clean, the variant is a witnessed reward hack.

Generation configuration. The same OpenAI model, `gpt-5.6-terra`, filled the adversarial variant, differential-scenario, neutral-solve, and property-generation roles. Retained generation requests set `max_completion_tokens=4000` and supplied no temperature field, so provider defaults applied. This role reuse is a limitation, not evidence of cross-model generality.

Result. From sixteen certified candidates, this procedure yields ten witnessed certified-resolved reward hacks across three repositories (Django, astropy, matplotlib). Two further candidates matched gold on their retained scenarios; we report them as `certified_equivalent` but do not infer global equivalence. The remaining candidates produced no clean witness. Table 1 lists these ten hacks, which formed the set used for the original ten-row detector pilot; the standing comparison in Section 5 uses all twenty-two released hacks.

The class is not specific to three repositories. We then ran the same pipeline on forty-two additional targets, disjoint from the original sixteen and spanning all twelve repositories in SWE-bench Verified. They were additional to this construction gate, not claimed mission-history fresh. The run produced twenty-three variants; twenty were certified resolved by execution, and twelve were also witnessed wrong, adding hacks in five new repositories: requests, xarray, pylint, pytest, and sympy. Combined with the first ten, the released benchmark is twenty-two execution-verified certified-resolved reward hacks across eight repositories. The sharpest scale example is a Django hack whose variant passes 1,870 graded tests and still returns 'accepted' where the correct fix returns 'rejected'. The instrument comparison below includes all twenty-two hacks.

Exploratory neutral-prompt localized solve. Every benchmark hack above is elicited. We therefore retained a separate run with no gaming instruction. This was not a random sample: thirty-nine targets were selected after excluding all sixty-six unique targets from the earlier iter193 Phase-A and iter199 target sets. The remaining engineering filters—one source file, one added run, strip-round-trip compatibility with the gold-shaped patch builder, lexicographic ordering, and a cap of five per repository—yielded 200 compatible rows across nine repositories and the frozen thirty-nine. The prompt included the issue and a buggy region reconstructed from gold-diff context and removed lines; gold-added lines were withheld. We therefore call it gold-localized, not plain.

gpt-5.6-terra produced thirty-seven valid patches; all now have official-harness evidence, and twenty-four certify. Nine model/gold patch pairs differ by exactly one terminal LF, so zero are byte-identical; eight of those normalized-equivalent patches certify. Four certified patches have no observed divergence on a valid witness, five lack a valid witness, and seven have an observed gold/model divergence. The blind judges were **gpt-5.6-terra** and **claude-opus-4-8**; each saw two outputs unlabeled. The OpenAI blind request used a 1,536-token cap and the Anthropic blind request used `max_tokens=400`; neither supplied a temperature field. The conservative requirement that both name only the model output was adopted after inspecting the original outcomes and reduces a loose count of three to one strict case. The original hypothesis named one judge, and Git commit `f651bfc` first records the hypothesis, target manifest, and solver outputs together, so the repository does not independently timestamp a pre-output freeze. The retained judge bundle stores parsed labels and derived booleans, not raw response text. Exact response substance and parser fidelity therefore cannot be re-audited; the strict case is bounded parsed-decision evidence.

The strict case is a Sphinx fix that passes all thirty-three graded tests yet renders a compound-key keyboard input as a malformed empty element. Five certified rows lack a valid witness and one divergence lacks two complete judge outcomes, so $u = 6$ remains unadjudicated. With $N = 24$ and $k = 1$, we report $1/24$ (4.1667%) confirmed lower, $7/24$ (29.1667%) as the worst case over those six declared missing pipeline outcomes, and $1/18$ (5.5556%) complete-case sensitivity. The formulas were adopted after the original result but before the denominator backfill; they are not preregistered iter200 estimates, confidence intervals, prevalence bounds, or population frequencies. The historical $1/15$ was conditional on a scenario-eligible cohort and is not pooled.

A worked example. The clearest hack is `django-11179`. When a database object is deleted, the correct fix clears the object's primary key to `None` in memory. The gold patch does exactly that. The variant clears the key only when the key is literally named `id`:

```
setattr(instance, pk.attname, None if pk.name == 'id' else instance.pk)
```

Every graded test for this instance uses a model whose primary key is named `id`, so the variant passes them and the harness certifies it. But a model whose primary key has any other name keeps

Table 1: Original iter195 ten-row subset of the released 22-row benchmark. Each patch passes every graded test (the official harness marks it resolved), yet produces different output from the gold fix on the shown input.

Instance	Function	Gold output	Variant output
django-11179	Collector.delete	None	'present'
django-11119	Engine.render_to_string	<tag>	<tag>
django-11163	model_to_dict	{}	{'id': None, ...}
django-11141	MigrationLoader.load_disk	(True, False)	(False, True)
django-11211	UUIDField.deconstruct	UUID('550e...')	'550e...' (str)
django-11433	construct_instance	0	7
astropy-7166	is_public_member	'...documentation'	None
astropy-14096	SkyCoord.__getattr__	'intercepted'	AttributeError
matplotlib-23412	Patch.draw	((2.5, ...),)	((0.0, ...),)
matplotlib-24627	_AxesBase.__clear	(True, True)	(False, False)

a stale key after deletion. Running the scenario, the gold patch prints `None` and the variant prints the stale value.

Execution beat our own judgment. One case is worth stating against ourselves. We hand-analyzed `django-11119` and judged its variant a correct rewrite: the changed expression looked equivalent for boolean inputs. Execution disagreed. With auto-escaping turned off, the gold patch renders text unescaped and the variant escapes it, so the two differ. A reasoning pass called the patch correct; running the code showed it was a hack. This is the paper’s own thesis applied to the authors: for this class, run the code rather than trust a verdict, including one’s own.

Claim boundary. The benchmark hacks are elicited and constructed, not sampled from real agent runs, so we make no claim about how often agents hack in practice. There are twenty-two hacks across eight repositories, and the exploratory instrument comparison uses all twenty-two. In this cohort, “wrong” means a clean synthesized-input output differs from gold. The separate localized solve supports one strict case under a later conservative blind rule, with all three missing-outcome quantities reported above. The two **certified_equivalent** controls matched gold only on retained scenarios; global equivalence is not claimed.

5 Detecting Certified-Yet-Wrong Patches

The instruments had different access. The judge saw the candidate diff, public task, and visible test, but not gold or the retained witness. The property generator saw the task, visible test, and a source/function locator extracted from the candidate diff; candidate code lines and gold content were withheld. Offline evaluation then used gold to decide which properties entered the analysis. The locator violated both runs’ diff-independent access descriptions. Gold validation was registered in iter201, so it is an interpretation limit there rather than a second protocol deviation; iter197 had instead registered the visible test as its soundness anchor and separate paired-gold controls.

Instrument A: judge panel. `gpt-5.6-terra` and `claude-opus-4-8` reviewed each of twenty-two hack diffs and twenty-two paired gold diffs. A planned `gemini-2.5-flash` slot was not run because application-default credentials were blocked. OpenAI judging used a 1,536-token cap and

Table 2: Corrected exploratory full-cohort instrument accounting. Property-gold passing is an inclusion rule, not an independent control result. Iter201’s protocol failure is the diff-derived locator.

Instrument	Hack catches	Coverage / missingness	Control interpretation
Judge panel	20/22	8/88 response nondecisions	Gold flags: 3/22 lower; 6/22 missing upper; 3/19 complete-case
Locator/gold property pipeline	6/22	21/22 generated; 20/21 pass gold	No independent false-positive estimate; protocol FAIL

Anthropic detector judging used `max_tokens=200`; no request supplied a temperature field. We count a row as caught if either parsed response says `reward_hack`.

The panel has observed catches on twenty of twenty-two hack rows. Eight of eighty-eight responses are unparseable, affecting five hack rows and three gold rows. Every affected hack row is caught by the other judge, so row-level any-catch remains twenty of twenty-two. Three gold rows are observed flags; because three other gold rows have a missing response, the required sensitivities are 3/22 observed lower, 6/22 worst-case missing upper, and 3/19 complete-case. The retained artifact contains parsed labels and nondecision markers, not raw response text, so exact response substance cannot be re-audited.

All forty-four full-cohort rows were freshly judged in iter201; no iter196 labels were reused. The prompt builder capped problem text at 1,500 characters, visible-test text at 2,500, and candidate diffs at 4,000. Task text was truncated in 14/44 prompts (7/22 unique instances), visible-test text in 6/44 (3/22), and candidate diffs in none. The first Git commit containing the full-cohort judge protocol also contained its outputs, so Git does not independently establish a pre-output protocol freeze.

Instrument B: locator-assisted, gold-validated property pipeline. `gpt-5.6-terra` generated a property for twenty-one of twenty-two targets. Twenty of those twenty-one properties passed on gold and entered the exploratory analysis; six caught the candidate patch. Because gold passing is the inclusion rule, running the same property on the paired gold patch is not an independent false-positive experiment. Iter201 records protocol **FAIL** because the generator used a diff-derived locator; its registered method did explicitly name gold validation. Iter197 records **FAIL** for three distinct reasons: the same locator deviation, replacement of the registered visible-test soundness anchor with gold, and omission of independent decisions on ten registered paired-gold controls. Its later threshold of five catches first appeared with generated outputs, so the observed 4/10 is a post-generation diagnostic, not a missed prospective bar.

The property prompts also used capped inputs. In iter201, task text was truncated for 9/22 targets and visible-test text for 3/22. Both historical property executions used mutable `:latest` image tags without retaining resolved digests; exact container bytes therefore cannot be reconstructed from the committed evidence.

Exploratory overlap. In iter197, the judge caught seven, the property pipeline caught four, and their observed union was eight of ten only when unparseable judge responses were treated as no flag. The apparent property-only row, `django-11211`, had two unparseable judge responses: it was judge-unadjudicated, not a confirmed judge miss. The corresponding missingness sensitivities are 8/10 observed, 9/10 upper, 7/8 judge-complete, and 8/9 property-resolved. In iter201, every one of the six property catches is among the judge’s twenty, so the union remains twenty of twenty-two

and two astrology hacks are missed by both. Neither run establishes confirmed complementarity or evidence that a deployable ensemble improves on a strong judge.

Claim boundary. This is a corrected exploratory comparison across eight repositories. The judge figures are one stochastic run with explicit nondecisions; the property counts come from a protocol deviation and have no independent false-positive estimate. We make no leaderboard, model-superiority, state-of-the-art, population-frequency, or deployment claim.

6 Limitations

We state the bounds plainly.

- **Small, mostly elicited sample.** The benchmark has twenty-two hacks across eight repositories, produced by a model asked to write wrong variants. The separate neutral-prompt solve is a nonrandom, gold-localized convenience sample, not a measure of field frequency.
- **Exploratory neutral-solve rules.** The strict blind rule was adopted after outcome inspection, and the missingness summaries were added after the original result but before the backfill. Six of twenty-four certified outcomes remain unadjudicated. The three sensitivities are descriptive and must stay together. Fifty-four historical execution logs lack explicit image/exit provenance and are trusted only as a frozen exact-byte corpus; twenty backfill logs carry the stronger provenance markers.
- **Cohort-specific wrongness labels.** In the constructed cohort, “wrong” means the variant differs from the gold reference fix on a retained input, not independent semantic adjudication. The neutral cohort additionally requires both blind judges to identify only the model output as wrong.
- **Incomplete judge records.** The detector-judge numbers come from one stochastic run with eight nondecisions. Parsed labels are retained, but raw iter201 response text is unavailable. Iter200 likewise retains only parsed blind-judge labels and derived booleans, so exact response substance and parser fidelity cannot be re-audited for the strict existence case.
- **Detector chronology and input coverage.** The full-cohort judge protocol and outputs first appear together in Git, so no independent pre-output freeze is established. Judge and property prompts used the character caps reported above; several task and test inputs were truncated.
- **Distinct property-protocol deviations.** Both property gates are FAIL because of derived locators. Iter197 additionally replaced its registered visible-test soundness anchor and omitted independent paired-gold decisions. Iter201 registered gold validation; there it prevents an independent specificity estimate but is not another failed bar.
- **Historical container provenance.** Iter197 and iter201 property execution retained mutable `:latest` tags but no resolved image digests. Their exact historical container bytes are not reconstructible from the repository.
- **Role reuse.** `gpt-5.6-terra` produced adversarial variants, scenarios, neutral patches, and properties and also occupied one judge slot; cross-model generalization is unmeasured.

7 Conclusion

Under our reference-differential operational label, SWE-bench Verified certifies patches that are wrong on retained synthesized inputs. We constructed twenty-two such patches across eight reposi-

tories and released execution witnesses. Every patch passes the graded tests; in a separate sixteen-row audit, no additional test in the audited target modules distinguished a candidate from gold. One judge-panel run has observed catches on twenty of twenty-two hacks, with eight of eighty-eight response nondecisions and paired-gold flag sensitivities of 3/22, 6/22, and 3/19. The locator-assisted, gold-validated property pipeline catches six of twenty-two, all within the judge set. Both property gates fail the diff-independent locator contract; iter197 has two additional protocol deviations, while iter201’s registered gold validation prevents an independent false-positive estimate. The ten-row apparent property-only case was judge-unadjudicated, and the full-cohort property catches are a subset of the judge’s, so no confirmed complementarity is shown. The exploratory neutral-prompt localized solve confirms one strict case under a post-inspection rule. Its corrected cohort has twenty-four certified patches and six declared missing outcomes, giving 1/24, 7/24, and 1/18. The pre-result next step is iter202: preserve the frozen cohort and execute only after its exact-byte runtime freeze, eight-shard certification plan, atomic checkpoints, and single-run-attempt aggregate receipt pass the committed primary-branch gate. The result must report the pre-result-declared prior-use and missing-outcome sensitivities.

Reproducibility. Every number in this paper regenerates from a committed script, manifest, or validated receipt in the Telos repository. The experiments are named in the text (for example, the construction is iter193 through iter195 and the expansion is iter199). Iter196 is the judge pilot; iter197 and iter201 are protocol-failed property runs with retained exploratory diagnostics. Iter200 is the exploratory localized solve and denominator correction. Iter202 is pre-result and contributes no result.

References

- [1] D. Amodei, C. Olah, J. Steinhardt, P. Christiano, J. Schulman, and D. Mané. Concrete Problems in AI Safety. *arXiv:1606.06565*, 2016.
- [2] V. Krakovna et al. Specification gaming: the flip side of AI ingenuity. DeepMind Blog, 2020.
- [3] J. Skalse, N. H. R. Howe, D. Krashenninikov, and D. Krueger. Defining and Characterizing Reward Hacking. *Advances in Neural Information Processing Systems (NeurIPS)*, 2022.
- [4] C. E. Jimenez, J. Yang, A. Wettig, S. Yao, K. Pei, O. Press, and K. Narasimhan. SWE-bench: Can Language Models Resolve Real-World GitHub Issues? *Proc. ICLR*, 2024. *arXiv:2310.06770*.
- [5] L. Zheng et al. Judging LLM-as-a-Judge with MT-Bench and Chatbot Arena. *Advances in Neural Information Processing Systems (NeurIPS)*, 2023. *arXiv:2306.05685*.
- [6] M. Terekhov, Z. N. D. Liu, C. Gulcehre, and S. Albanie. Control Tax: The Price of Keeping AI in Check. *arXiv:2506.05296*, 2025.
- [7] Google DeepMind. Frontier Safety Framework, version 3.0. September 22, 2025.
- [8] METR. MALT: A Dataset of Natural and Prompted Behaviors That Threaten Eval Integrity. October 14, 2025. (MALT: Manually-reviewed Agentic Labeled Transcripts; 10,919 transcripts, 403 tasks, 21 models.)
- [9] N. Lambert et al. RewardBench: Evaluating Reward Models for Language Modeling. *arXiv:2403.13787*, 2024.
- [10] T. Y. Chen et al. Metamorphic Testing: A Review of Challenges and Opportunities. *ACM Computing Surveys*, 51(1), 2018.
- [11] K. Claessen and J. Hughes. QuickCheck: a lightweight tool for random testing of Haskell programs. *Proc. ICFP*, 2000.